

Assignment 5.1 - The IMDB Dataset

Loading Data Set

```
In [1]: import numpy as np
        from tensorflow.keras.datasets import imdb
```

```
In [2]: # splitting into training and test dataset.
        # each review is list of word indices. train labels: 0 = negative, 1 = positive.
        (train_data, train_labels), (test_data, test_labels) = imdb.load_data(num_words=10000)
```

```
In [3]: max([max(sequence) for sequence in train_data])
```

Out[3]: 9999

```
In [4]: # maps words to an integer index
        word_index = imdb.get_word_index()
        #word_index
```

```
In [5]: # maps integer indices to words
        reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])
        decoded_review = ' '.join([reverse_word_index.get(i - 3, '?') for i in train_data[0]])
        decoded_review
```

Out[5]: "? this film was just brilliant casting location scenery story direction everyone's really suited the part they played and you could just imagine being there robert ? is an amazing actor and now the same being director ? father came from the same scottish island as myself so i loved the fact there was a real connection with this film the witty remarks throughout the film were great it was just brilliant so much that i bought the film as soon as it was released for ? and would recommend it to everyone to watch and the fly fishing was amazing really cried at the end it was so sad and you know what they say if you cry at a film it must have been good and this definitely was also ? to the two little boy's that played the ? of norman and paul they were just brilliant children are often left out of the ? list i think because the stars that play them all grown up are such a big profile for the whole film but these children are amazing and should be praised for what they have done don't you think the whole story was so lovely because it was true and was someone's life after all that was shared with us all"

Preparing Dataset

```
In [6]: def vectorize_sequences(sequences, dimension=10000):  
        results = np.zeros((len(sequences), dimension))  
        for i, sequence in enumerate(sequences):  
            results[i, sequence] = 1  
        return results  
  
x_train = vectorize_sequences(train_data)  
x_test = vectorize_sequences(test_data)
```

```
In [7]: x_train[0]
```

```
Out[7]: array([0., 1., 1., ..., 0., 0., 0.])
```

```
In [8]: y_train = np.asarray(train_labels).astype('float32')  
y_test = np.asarray(test_labels).astype('float32')
```

Building your network

```
In [9]: # Import keras libraries  
from keras import models, layers, losses, metrics, optimizers
```

Using TensorFlow backend.

```
In [10]: # Creating a neural net with shape (1000,16,16,1)  
model = models.Sequential()  
model.add(layers.Dense(16, activation = 'relu', input_shape = (10000,)))  
model.add(layers.Dense(16, activation = 'relu'))  
model.add(layers.Dense(1, activation = 'sigmoid'))
```

```
In [11]: model.compile(optimizer='rmsprop',  
                      loss = 'binary_crossentropy',  
                      metrics = ['accuracy'])
```

```
In [12]: model.compile(optimizer = optimizers.RMSprop(lr = 0.001),  
                      loss = 'binary_crossentropy',  
                      metrics = ['accuracy'])
```

```
In [13]: # using custom losses and metrics  
model.compile(optimizer=optimizers.RMSprop(lr=0.001),  
              loss=losses.binary_crossentropy,  
              metrics=[metrics.binary_accuracy])
```

Validating the model

```
In [14]: # Splitting the data into validation and train set
x_val = x_train[:10000]
partial_x_train = x_train[10000:]

y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

Training the model

```
In [15]: model.compile(optimizer='rmsprop',
                        loss = 'binary_crossentropy',
                        metrics = ['acc'])
```

```
In [16]: # Training the neural network with partial_x_train and partial_y_train
history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs = 20,
                    batch_size = 512,
                    validation_data = (x_val, y_val))
```

Train on 15000 samples, validate on 10000 samples

Epoch 1/20

15000/15000 [=====] - 3s 197us/step - loss: 0.5084 -
acc: 0.7813 - val_loss: 0.3797 - val_acc: 0.8684

Epoch 2/20

15000/15000 [=====] - 2s 153us/step - loss: 0.3004 -
acc: 0.9047 - val_loss: 0.3004 - val_acc: 0.8897

Epoch 3/20

15000/15000 [=====] - 2s 152us/step - loss: 0.2179 -
acc: 0.9285 - val_loss: 0.3085 - val_acc: 0.8711

Epoch 4/20

15000/15000 [=====] - 2s 157us/step - loss: 0.1750 -
acc: 0.9437 - val_loss: 0.2840 - val_acc: 0.8832

Epoch 5/20

15000/15000 [=====] - 2s 145us/step - loss: 0.1427 -
acc: 0.9543 - val_loss: 0.2841 - val_acc: 0.8872

Epoch 6/20

15000/15000 [=====] - 2s 149us/step - loss: 0.1150 -
acc: 0.9650 - val_loss: 0.3166 - val_acc: 0.8772

Epoch 7/20

15000/15000 [=====] - 2s 152us/step - loss: 0.0980 -
acc: 0.9705 - val_loss: 0.3127 - val_acc: 0.8846

Epoch 8/20

15000/15000 [=====] - 2s 155us/step - loss: 0.0807 -
acc: 0.9763 - val_loss: 0.3861 - val_acc: 0.8649

Epoch 9/20

15000/15000 [=====] - 2s 154us/step - loss: 0.0661 -
acc: 0.9821 - val_loss: 0.3637 - val_acc: 0.8782

Epoch 10/20

15000/15000 [=====] - 2s 148us/step - loss: 0.0561 -
acc: 0.9853 - val_loss: 0.3850 - val_acc: 0.8792

Epoch 11/20

15000/15000 [=====] - 2s 155us/step - loss: 0.0439 -
acc: 0.9893 - val_loss: 0.4168 - val_acc: 0.8779

Epoch 12/20

15000/15000 [=====] - 2s 154us/step - loss: 0.0381 -
acc: 0.9921 - val_loss: 0.4552 - val_acc: 0.8690

Epoch 13/20

15000/15000 [=====] - 2s 152us/step - loss: 0.0300 -
acc: 0.9928 - val_loss: 0.4733 - val_acc: 0.8729

Epoch 14/20

15000/15000 [=====] - 2s 157us/step - loss: 0.0247 -
acc: 0.9945 - val_loss: 0.5068 - val_acc: 0.8726

Epoch 15/20

15000/15000 [=====] - 2s 153us/step - loss: 0.0175 -
acc: 0.9979 - val_loss: 0.5412 - val_acc: 0.8693

Epoch 16/20

15000/15000 [=====] - 2s 153us/step - loss: 0.0149 -
acc: 0.9983 - val_loss: 0.5802 - val_acc: 0.8699

Epoch 17/20

15000/15000 [=====] - 2s 155us/step - loss: 0.0151 -
acc: 0.9971 - val_loss: 0.6151 - val_acc: 0.8697

Epoch 18/20

15000/15000 [=====] - 2s 156us/step - loss: 0.0075 -
acc: 0.9996 - val_loss: 0.6948 - val_acc: 0.8632

Epoch 19/20

15000/15000 [=====] - 2s 155us/step - loss: 0.0121 -

```
acc: 0.9971 - val_loss: 0.6941 - val_acc: 0.8652
Epoch 20/20
15000/15000 [=====] - 2s 152us/step - loss: 0.0041 -
acc: 0.9999 - val_loss: 0.7117 - val_acc: 0.8657
```

```
In [17]: history_dict = history.history
         history_dict.keys()
```

```
Out[17]: dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```


Plotting the training and validation loss

```
In [18]: import matplotlib.pyplot as plt
         acc = history_dict['acc']
         val_acc = history_dict['val_acc']


         loss_values = history_dict['loss']
         val_loss_values = history_dict['val_loss']

         epochs = range(1, len(acc) + 1)

         plt.plot(epochs, loss_values, 'bo', label = 'Training Loss')
         plt.plot(epochs, val_loss_values, 'b', label = 'Validation loss')
         plt.title('Training and Validation Loss')
         plt.xlabel("Epochs")
         plt.ylabel("Loss")
         plt.legend()
         plt.show()
```



```
In [19]: plt.plot(epochs, acc, 'bo', label = 'Training accuracy')
plt.plot(epochs, val_acc, 'b', label = 'Validation accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



```
In [20]: # Retraining a model from scratch
model = models.Sequential()
model.add(layers.Dense(16, activation = 'relu', input_shape = (10000,)))
model.add(layers.Dense(16, activation = 'relu'))
model.add(layers.Dense(1, activation = 'sigmoid'))

model.compile(optimizer='rmsprop',
              loss = 'binary_crossentropy',
              metrics = ['accuracy'])

model.fit(partial_x_train,
          partial_y_train,
          epochs = 4,
          batch_size = 512)

results = model.evaluate(x_test, y_test)
results
```

```
Epoch 1/4
15000/15000 [=====] - 2s 109us/step - loss: 0.5326 -
accuracy: 0.7917
Epoch 2/4
15000/15000 [=====] - 2s 101us/step - loss: 0.3258 -
accuracy: 0.8987
Epoch 3/4
15000/15000 [=====] - 1s 97us/step - loss: 0.2357 -
accuracy: 0.9245
Epoch 4/4
15000/15000 [=====] - 2s 102us/step - loss: 0.1866 -
accuracy: 0.9397
25000/25000 [=====] - 4s 156us/step
```

```
Out[20]: [0.30134918537139893, 0.8772000074386597]
```

Use trained model to generate predictions

```
In [21]: model.predict(x_test)
```

```
Out[21]: array([[0.33491266],
                [0.9995926 ],
                [0.9345059 ],
                ...,
                [0.14445111],
                [0.19112667],
                [0.65519667]], dtype=float32)
```

```
In [ ]:
```