

## Assignment 6.2a

```
In [1]: import keras
import tensorflow as tf
from keras.datasets import cifar10

from keras import layers
from keras import models
import matplotlib.pyplot as plt
import numpy as np
from keras.preprocessing.image import ImageDataGenerator
```

Using TensorFlow backend.

### Load Data Set

```
In [2]: # Loading the data into memory
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

```
In [3]: x_train.shape
```

```
Out[3]: (50000, 32, 32, 3)
```

```
In [4]: y_train.shape
```

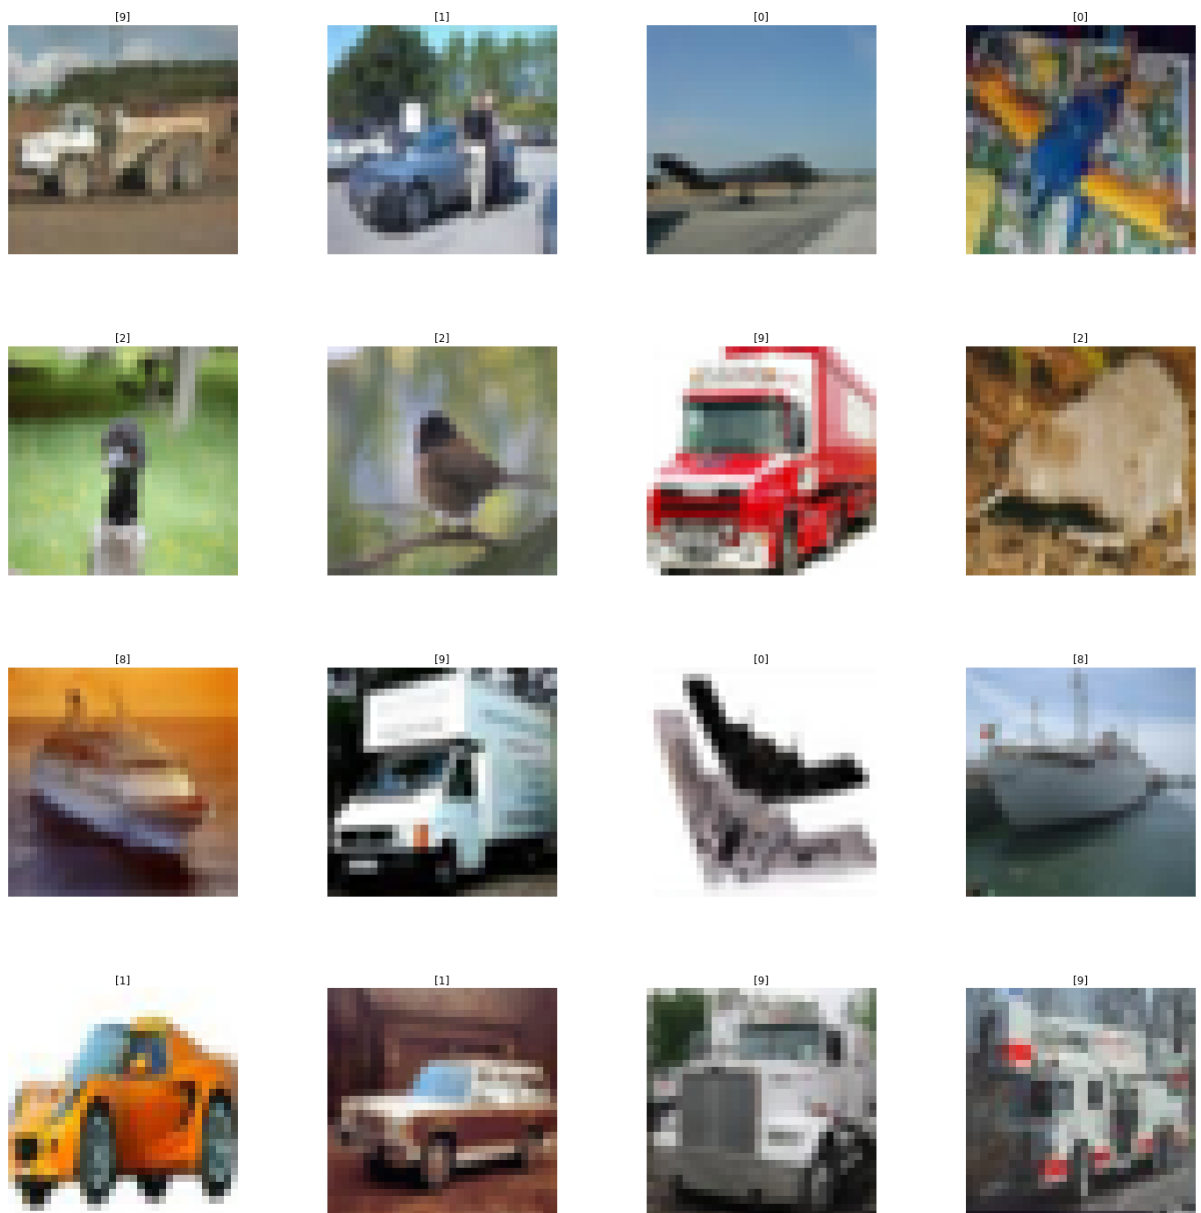
```
Out[4]: (50000, 1)
```

```

In [5]: W_grid = 4
        L_grid = 4
        fig, axes = plt.subplots(L_grid, W_grid, figsize = (25, 25))
        axes = axes.ravel()
        n_training = len(x_train)
        for i in np.arange(0, L_grid * W_grid):
            index = np.random.randint(0, n_training) # pick a random number
            axes[i].imshow(x_train[index])
            axes[i].set_title(y_train[index])
            axes[i].axis('off')
        plt.subplots_adjust(hspace = 0.4)

```

C:\Users\bibek\anaconda3\envs\dsc650\lib\site-packages\matplotlib\text.py:115  
 0: FutureWarning: elementwise comparison failed; returning scalar instead, but  
 in the future will perform elementwise comparison  
 if s != self.\_text:



## Data Preparation

```
In [6]: x_train = x_train.astype('float32') / 255
        x_test = x_test.astype('float32') / 255

        y_train = y_train.reshape(y_train.shape[0])
        y_test = y_test.reshape(y_test.shape[0])
```

### Validation Dataset

```
In [7]: x_val_train = x_train[:10000]
        x_train = x_train[10000:]

        y_val_train = y_train[:10000]
        y_train = y_train[10000:]
```

### Instantiating a Small Covnet

```
In [8]: model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', padding = 'same',
                        input_shape=x_train.shape[1:]))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(10, activation='sigmoid'))

model.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 14, 14, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_3 (Conv2D)	(None, 5, 5, 128)	73856
max_pooling2d_3 (MaxPooling2D)	(None, 2, 2, 128)	0
flatten_1 (Flatten)	(None, 512)	0
dense_1 (Dense)	(None, 512)	262656
dense_2 (Dense)	(None, 10)	5130
=====		
Total params: 361,034		
Trainable params: 361,034		
Non-trainable params: 0		

```
In [9]: from keras import optimizers
```

## Building a Neural Network

```
In [10]: model.compile(optimizer=optimizers.RMSprop(lr=1e-4),  
                        loss='sparse_categorical_crossentropy',  
                        metrics=['accuracy'])  
  
# Training the neural network with partial_x_train and partial_y_train  
history = model.fit(x_train,  
                    y_train,  
                    epochs = 20,  
                    batch_size = 128,  
                    validation_data = (x_val_train, y_val_train),  
                    verbose=True)
```

Train on 40000 samples, validate on 10000 samples

Epoch 1/20

40000/40000 [=====] - 21s 535us/step - loss: 2.0406  
- accuracy: 0.2572 - val\_loss: 1.9585 - val\_accuracy: 0.2711

Epoch 2/20

40000/40000 [=====] - 22s 555us/step - loss: 1.7731  
- accuracy: 0.3532 - val\_loss: 1.7868 - val\_accuracy: 0.3545

Epoch 3/20

40000/40000 [=====] - 23s 568us/step - loss: 1.6595  
- accuracy: 0.3988 - val\_loss: 1.6384 - val\_accuracy: 0.3896

Epoch 4/20

40000/40000 [=====] - 23s 565us/step - loss: 1.5841  
- accuracy: 0.4273 - val\_loss: 1.6020 - val\_accuracy: 0.4266

Epoch 5/20

40000/40000 [=====] - 23s 571us/step - loss: 1.5263  
- accuracy: 0.4488 - val\_loss: 1.4835 - val\_accuracy: 0.4588

Epoch 6/20

40000/40000 [=====] - 23s 567us/step - loss: 1.4794  
- accuracy: 0.4675 - val\_loss: 1.5183 - val\_accuracy: 0.4457

Epoch 7/20

40000/40000 [=====] - 23s 571us/step - loss: 1.4357  
- accuracy: 0.4892 - val\_loss: 1.4310 - val\_accuracy: 0.4912

Epoch 8/20

40000/40000 [=====] - 23s 573us/step - loss: 1.3988  
- accuracy: 0.5029 - val\_loss: 1.3958 - val\_accuracy: 0.4993

Epoch 9/20

40000/40000 [=====] - 23s 574us/step - loss: 1.3655  
- accuracy: 0.5174 - val\_loss: 1.3653 - val\_accuracy: 0.5070

Epoch 10/20

40000/40000 [=====] - 23s 576us/step - loss: 1.3389  
- accuracy: 0.5298 - val\_loss: 1.3489 - val\_accuracy: 0.5270

Epoch 11/20

40000/40000 [=====] - 23s 576us/step - loss: 1.3114  
- accuracy: 0.5385 - val\_loss: 1.2929 - val\_accuracy: 0.5456

Epoch 12/20

40000/40000 [=====] - 24s 593us/step - loss: 1.2879  
- accuracy: 0.5466 - val\_loss: 1.4090 - val\_accuracy: 0.5040

Epoch 13/20

40000/40000 [=====] - 24s 592us/step - loss: 1.2605  
- accuracy: 0.5585 - val\_loss: 1.3054 - val\_accuracy: 0.5334

Epoch 14/20

40000/40000 [=====] - 23s 579us/step - loss: 1.2408  
- accuracy: 0.5646 - val\_loss: 1.3223 - val\_accuracy: 0.5295

Epoch 15/20

40000/40000 [=====] - 24s 588us/step - loss: 1.2183  
- accuracy: 0.5738 - val\_loss: 1.2371 - val\_accuracy: 0.5611

Epoch 16/20

40000/40000 [=====] - 25s 636us/step - loss: 1.1993  
- accuracy: 0.5814 - val\_loss: 1.2090 - val\_accuracy: 0.5710

Epoch 17/20

40000/40000 [=====] - 27s 684us/step - loss: 1.1812  
- accuracy: 0.5889 - val\_loss: 1.2005 - val\_accuracy: 0.5681

Epoch 18/20

40000/40000 [=====] - 27s 677us/step - loss: 1.1643  
- accuracy: 0.5945 - val\_loss: 1.1959 - val\_accuracy: 0.5809

Epoch 19/20

40000/40000 [=====] - 24s 596us/step - loss: 1.1461

```
- accuracy: 0.6010 - val_loss: 1.2313 - val_accuracy: 0.5581
Epoch 20/20
40000/40000 [=====] - 28s 703us/step - loss: 1.1307
- accuracy: 0.6061 - val_loss: 1.1454 - val_accuracy: 0.5972
```

```
In [11]: results = model.evaluate(x_test, y_test)
```

```
results
```

```
10000/10000 [=====] - 2s 218us/step
```

```
Out[11]: [1.1482285236358643, 0.597699998092651]
```

```
In [12]: history_dict = history.history
history_dict.keys()
```


```
Out[12]: dict_keys(['val_loss', 'val_accuracy', 'loss', 'accuracy'])
```

```
In [13]: acc = history_dict['accuracy']
val_acc = history_dict['val_accuracy']


loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']

epochs = range(1, len(acc) + 1)

plt.plot(epochs, loss_values, "r-", label = 'Training Loss')
plt.plot(epochs, val_loss_values, label = 'Validation loss')
plt.title('Training and Validation Loss')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



```
In [14]: plt.plot(epochs, acc, "r-", label = 'Training accuracy')
plt.plot(epochs, val_acc, label = 'Validation accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



## Assignment 6.2b



```
In [15]: model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', padding = 'same',
                        input_shape=x_train.shape[1:]))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(10, activation='sigmoid'))

model.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_4 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_5 (Conv2D)	(None, 14, 14, 64)	18496
max_pooling2d_5 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_6 (Conv2D)	(None, 5, 5, 128)	73856
max_pooling2d_6 (MaxPooling2D)	(None, 2, 2, 128)	0
flatten_2 (Flatten)	(None, 512)	0
dropout_1 (Dropout)	(None, 512)	0
dense_3 (Dense)	(None, 512)	262656
dense_4 (Dense)	(None, 10)	5130
=====		
Total params: 361,034		
Trainable params: 361,034		
Non-trainable params: 0		
=====		

```
In [16]: model.compile(optimizer=optimizers.RMSprop(lr=1e-4),
                    loss='sparse_categorical_crossentropy',
                    metrics=['accuracy'])

train_datagen = ImageDataGenerator(
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True)

train_datagen.fit(x_train)

# Training the neural network with partial_x_train and partial_y_train
history = model.fit(train_datagen.flow(x_train, y_train,
                                       batch_size=128), epochs=40,
                   validation_data=(x_val_train, y_val_train),
                   workers=4, verbose=False)
```

```
In [17]: results = model.evaluate(x_test, y_test)
```

```
results
```

```
10000/10000 [=====] - 3s 255us/step
```

```
Out[17]: [1.2589075031280517, 0.5558000206947327]
```

```
In [18]: history_dict = history.history
        history_dict.keys()
```


```
Out[18]: dict_keys(['val_loss', 'val_accuracy', 'loss', 'accuracy'])
```

```
In [19]: acc = history_dict['accuracy']
val_acc = history_dict['val_accuracy']


loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']

epochs = range(1, len(acc) + 1)

plt.plot(epochs, loss_values, "r-", label = 'Training Loss')
plt.plot(epochs, val_loss_values, label = 'Validation loss')
plt.title('Training and Validation Loss')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



```
In [20]: plt.plot(epochs, acc, "r-", label = 'Training accuracy')
plt.plot(epochs, val_acc, label = 'Validation accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



In [ ]: