

Assignment 10

DSCT650

Taniya Adhikari

Assignment 10.1a

```
In [1]: import re
        from string import punctuation
```

```
In [2]: def tokenize(sentence, split = " "):

        # removes punctuation
        pattern = r'^A-Za-z ]'
        regex = re.compile(pattern)
        sentence = regex.sub(' ', sentence)
        sentence = sentence.replace("\n", "")

        # Convert string to lower case
        sentence = sentence.lower()

        #removes duplicate spaces
        sentence = re.sub(' +', ' ', sentence)

        # remove whitespace
        sentence = sentence.strip()

        # Split by spaces
        tokens = sentence.split(split)
        return tokens
```

```
In [3]: s = """
        She stared at it, and rubbed her eyes, and stared at it again.

        "Well! I never!" she said at last. "And me thinking it was a pot of gold!
        I must have been dreaming. But this is luck! Silver is far less trouble—easier
        to mind, and not so easy stolen. Them gold pieces would have been the death o'
        me,
        and with this great lump of silver—"

        """
```

```
In [4]: tokens = tokenize(s)
print(tokens)
```

```
['she', 'stared', 'at', 'it', 'and', 'rubbed', 'her', 'eyes', 'and', 'stare
d', 'at', 'it', 'again', 'well', 'i', 'never', 'she', 'said', 'at', 'last',
'and', 'me', 'thinking', 'it', 'was', 'a', 'pot', 'of', 'gold', 'i', 'must',
'have', 'been', 'dreaming', 'but', 'this', 'is', 'luck', 'silver', 'is', 'fa
r', 'less', 'trouble', 'easier', 'to', 'mind', 'and', 'not', 'so', 'easy', 's
tolen', 'them', 'gold', 'pieces', 'would', 'have', 'been', 'the', 'death',
'o', 'me', 'and', 'with', 'this', 'great', 'lump', 'of', 'silver']
```

Assignment 10.1b

```
In [5]: def ngram(tokens, n):
        ngrams = []

        for num in range(0, len(tokens)):
            ngram = ' '.join(tokens[num:num + n])
            ngrams.append(ngram)

        return ngrams

print(ngram(tokens, 2))
```

```
['she stared', 'stared at', 'at it', 'it and', 'and rubbed', 'rubbed her', 'h
er eyes', 'eyes and', 'and stared', 'stared at', 'at it', 'it again', 'again
well', 'well i', 'i never', 'never she', 'she said', 'said at', 'at last', 'l
ast and', 'and me', 'me thinking', 'thinking it', 'it was', 'was a', 'a pot',
'pot of', 'of gold', 'gold i', 'i must', 'must have', 'have been', 'been drea
ming', 'dreaming but', 'but this', 'this is', 'is luck', 'luck silver', 'silv
er is', 'is far', 'far less', 'less trouble', 'trouble easier', 'easier to',
'to mind', 'mind and', 'and not', 'not so', 'so easy', 'easy stolen', 'stolen
them', 'them gold', 'gold pieces', 'pieces would', 'would have', 'have been',
'been the', 'the death', 'death o', 'o me', 'me and', 'and with', 'with thi
s', 'this great', 'great lump', 'lump of', 'of silver', 'silver']
```

Assignment 10.1c

```
In [6]: import numpy as np
def one_hot_encode(tokens, num_words):
    token_index = {}
    for word in tokens:
        if word not in token_index:
            token_index[word] = len(token_index) + 1

    results = np.zeros(shape=(num_words, max(token_index.values()) + 1))
    for j, word in enumerate(tokens):
        index = token_index.get(word)
        results[j, index] = 1.
    return results
```

```
In [7]: one_hot_encode(tokens, len(tokens))
```

```
Out[7]: array([[0., 1., 0., ..., 0., 0., 0.],
               [0., 0., 1., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               ...,
               [0., 0., 0., ..., 0., 0., 1.],
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.]])
```

Assignment 10.2

```
In [8]: import pandas as pd
import matplotlib.pyplot as plt
import tensorflow

from tensorflow import keras
from keras import models, layers, losses, optimizers
from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
```

Using TensorFlow backend.

```
In [9]: import os
from pathlib import Path
import shutil, random
```

```
In [10]: current_dir = Path(os.getcwd()).absolute()
current_dir = Path(current_dir).parents[2]
current_dir
```

```
Out[10]: WindowsPath('C:/Users/bibek/Documents/GitHub/dsc650')
```

```
In [11]: data_dir = current_dir.joinpath('data')
external_dir = data_dir.joinpath('external')
imdb_dir = external_dir.joinpath('imdb')
base_dir = imdb_dir.joinpath("aclImdb")
base_dir
```

```
Out[11]: WindowsPath('C:/Users/bibek/Documents/GitHub/dsc650/data/external/imdb/aclImdb')
```

```
In [12]: train_dir = base_dir.joinpath("train")
test_dir = base_dir.joinpath("test")
```

```
In [13]: def data_set(directory):
    labels = []
    texts = []
    for label_type in ['neg', 'pos']:
        dir_name = os.path.join(directory, label_type)
        for root, dirs, files in os.walk(dir_name):
            for file in files:
                current_path = Path(root).joinpath(file)
                with open(current_path, encoding="utf8") as f:
                    review = f.read()
                    texts.append(review)
                    f.close()
                if label_type == 'neg':
                    labels.append(0)
                else:
                    labels.append(1)
    return texts, labels
```

```
In [14]: X_train, Y_train = data_set(train_dir)
    X_test, Y_test = data_set(test_dir)
```

```
In [15]: X_train[0], Y_train[0]
```

```
Out[15]: ("Story of a man who has unnatural feelings for a pig. Starts out with a open
ing scene that is a terrific example of absurd comedy. A formal orchestra aud
ience is turned into an insane, violent mob by the crazy chantings of it's si
ngers. Unfortunately it stays absurd the WHOLE time with no general narrative
eventually making it just too off putting. Even those from the era should be
turned off. The cryptic dialogue would make Shakespeare seem easy to a third
grader. On a technical level it's better than you might think with some good
cinematography by future great Vilmos Zsigmond. Future stars Sally Kirkland a
nd Frederic Forrest can be seen briefly.",
0)
```

```
In [16]: X_test[0], Y_test[0]
```

```
Out[16]: ("Once again Mr. Costner has dragged out a movie for far longer than necessar
y. Aside from the terrific sea rescue sequences, of which there are very few
I just did not care about any of the characters. Most of us have ghosts in th
e closet, and Costner's character are realized early on, and then forgotten u
ntil much later, by which time I did not care. The character we should really
care about is a very cocky, overconfident Ashton Kutcher. The problem is he c
omes off as kid who thinks he's better than anyone else around him and shows
no signs of a cluttered closet. His only obstacle appears to be winning over
Costner. Finally when we are well past the half way point of this stinker, Co
stner tells us all about Kutcher's ghosts. We are told why Kutcher is driven
to be the best with no prior inkling or foreshadowing. No magic here, it was
all I could do to keep from turning it off an hour in.",
0)
```

```
In [17]: def preprocess_data(review):  
    # removes punctuation  
    pattern = r'^A-Za-z ]'  
    regex = re.compile(pattern)  
    review = regex.sub(' ', review)  
    review = review.replace("\n", "")  
  
    # Convert string to Lower case  
    review = review.lower()  
    #removes duplicate spaces  
    review = re.sub(' +', ' ', review)  
    # remove whitespace  
    review = review.strip()  
  
    return review
```

```
In [18]: X_train = [preprocess_data(review) for review in X_train]  
X_test = [preprocess_data(review) for review in X_test]
```

```
In [19]: len(X_train)
```

```
Out[19]: 25000
```

Tokenizing

```
In [20]: max_words = 10000  
tokenizer = Tokenizer(num_words = max_words)  
tokenizer.fit_on_texts(X_train)  
sequences = tokenizer.texts_to_sequences(X_train)  
word_index = tokenizer.word_index  
print('Found %s unique tokens.' % len(word_index))
```

Found 73272 unique tokens.

```
In [21]: max_len = max([len(x) for x in sequences])  
data = pad_sequences(sequences, maxlen=max_len)  
labels = np.asarray(Y_train)  
  
print('Shape of data tensor:', data.shape)  
print('Shape of label tensor:', labels.shape)
```

Shape of data tensor: (25000, 2234)
Shape of label tensor: (25000,)

```
In [22]: indices = np.arange(data.shape[0])  
np.random.shuffle(indices)  
data = data[indices]  
labels = labels[indices]
```

```
In [23]: # Split the data into training and validation sets
X_train = data[500:]
Y_train = labels[500:]

X_val = data[:500]
Y_val = labels[:500]
```

Training the Model

```
In [24]: embedding_dim = 100

model = Sequential()
model.add(Embedding(max_words,
                    embedding_dim,
                    input_length=max_len))
model.add(Flatten())
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
embedding_1 (Embedding)	(None, 2234, 100)	1000000
=====		
flatten_1 (Flatten)	(None, 223400)	0
=====		
dense_1 (Dense)	(None, 32)	7148832
=====		
dense_2 (Dense)	(None, 1)	33
=====		
Total params: 8,148,865		
Trainable params: 8,148,865		
Non-trainable params: 0		
=====		

```
In [25]: model.compile(optimizer='rmsprop',
                      loss='binary_crossentropy',
                      metrics=['acc'])

history = model.fit(X_train,Y_train,
                   epochs=10,
                   batch_size=32,
                   validation_data=(X_val, Y_val))
```

C:\Users\bibek\anaconda3\envs\dsc650\lib\site-packages\tensorflow_core\python\framework\indexed_slices.py:433: UserWarning: Converting sparse IndexedSlices to a dense Tensor of unknown shape. This may consume a large amount of memory.

"Converting sparse IndexedSlices to a dense Tensor of unknown shape. "

Train on 24500 samples, validate on 500 samples

Epoch 1/10

24500/24500 [=====] - 51s 2ms/step - loss: 0.4495 - acc: 0.7946 - val_loss: 0.3443 - val_acc: 0.8600

Epoch 2/10

24500/24500 [=====] - 53s 2ms/step - loss: 0.1559 - acc: 0.9409 - val_loss: 0.3847 - val_acc: 0.8720

Epoch 3/10

24500/24500 [=====] - 55s 2ms/step - loss: 0.0419 - acc: 0.9859 - val_loss: 0.5166 - val_acc: 0.8620

Epoch 4/10

24500/24500 [=====] - 54s 2ms/step - loss: 0.0125 - acc: 0.9959 - val_loss: 0.7283 - val_acc: 0.8760

Epoch 5/10

24500/24500 [=====] - 52s 2ms/step - loss: 0.0169 - acc: 0.9962 - val_loss: 0.9751 - val_acc: 0.8400

Epoch 6/10

24500/24500 [=====] - 52s 2ms/step - loss: 0.0090 - acc: 0.9985 - val_loss: 1.1550 - val_acc: 0.8600

Epoch 7/10

24500/24500 [=====] - 52s 2ms/step - loss: 0.0104 - acc: 0.9980 - val_loss: 1.2676 - val_acc: 0.8500

Epoch 8/10

24500/24500 [=====] - 52s 2ms/step - loss: 0.0036 - acc: 0.9991 - val_loss: 1.3762 - val_acc: 0.8600

Epoch 9/10

24500/24500 [=====] - 52s 2ms/step - loss: 0.0114 - acc: 0.9986 - val_loss: 1.5495 - val_acc: 0.8500

Epoch 10/10

24500/24500 [=====] - 52s 2ms/step - loss: 0.0016 - acc: 0.9996 - val_loss: 1.7535 - val_acc: 0.8280

Evaluating the Model

```
In [26]: X_test = pad_sequences(sequences, maxlen=max_len)
         Y_test = np.asarray(labels)
```

In [27]: `model.evaluate(X_test,Y_test)`

25000/25000 [=====] - 7s 279us/step

Out[27]: [13.71816395401001, 0.4968000054359436]


```
In [28]: history_dict = history.history


acc = history_dict['acc']
val_acc = history_dict['val_acc']
loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']
epochs = range(1, len(acc) + 1)

# Plotting metrics
plt.plot(epochs, acc, 'bo', label = 'Training accuracy')
plt.plot(epochs, val_acc, 'b', label = 'Validation accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()

plt.figure()

plt.plot(epochs, loss_values, 'bo', label = 'Training Loss')
plt.plot(epochs, val_loss_values, 'b', label = 'Validation loss')
plt.title('Training and Validation Loss')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()

plt.show()
```



Assignment 10.3

```
In [29]: from keras.layers import LSTM
```

```
In [30]: max_features = 10000
model = Sequential()
model.add(Embedding(max_words,
                    embedding_dim,
                    input_length=max_len))
model.add(LSTM(32))
model.add(Dense(1, activation='sigmoid'))
```

```
In [31]: model.compile(optimizer='rmsprop',
                      loss='binary_crossentropy',
                      metrics=['acc'])
history = model.fit(X_train, Y_train,
                   epochs=10,
                   batch_size=128,
                   validation_data=(X_val, Y_val))
```

C:\Users\bibek\anaconda3\envs\dsc650\lib\site-packages\tensorflow_core\python\framework\indexed_slices.py:433: UserWarning: Converting sparse IndexedSlices to a dense Tensor of unknown shape. This may consume a large amount of memory.

"Converting sparse IndexedSlices to a dense Tensor of unknown shape. "

Train on 24500 samples, validate on 500 samples

Epoch 1/10

24500/24500 [=====] - 495s 20ms/step - loss: 0.4544
- acc: 0.7927 - val_loss: 0.3596 - val_acc: 0.8620

Epoch 2/10

24500/24500 [=====] - 466s 19ms/step - loss: 0.2797
- acc: 0.8907 - val_loss: 0.4469 - val_acc: 0.8580

Epoch 3/10

24500/24500 [=====] - 470s 19ms/step - loss: 0.2318
- acc: 0.9129 - val_loss: 0.3020 - val_acc: 0.8900

Epoch 4/10

24500/24500 [=====] - 449s 18ms/step - loss: 0.2031
- acc: 0.9251 - val_loss: 0.3877 - val_acc: 0.8680

Epoch 5/10

24500/24500 [=====] - 474s 19ms/step - loss: 0.1830
- acc: 0.9338 - val_loss: 0.3534 - val_acc: 0.8840

Epoch 6/10

24500/24500 [=====] - 474s 19ms/step - loss: 0.1638
- acc: 0.9401 - val_loss: 0.4127 - val_acc: 0.8740

Epoch 7/10

24500/24500 [=====] - 505s 21ms/step - loss: 0.1491
- acc: 0.9473 - val_loss: 0.4494 - val_acc: 0.8680

Epoch 8/10

24500/24500 [=====] - 506s 21ms/step - loss: 0.1447
- acc: 0.9498 - val_loss: 0.4349 - val_acc: 0.8740

Epoch 9/10

24500/24500 [=====] - 467s 19ms/step - loss: 0.1198
- acc: 0.9573 - val_loss: 0.4533 - val_acc: 0.8640

Epoch 10/10

24500/24500 [=====] - 514s 21ms/step - loss: 0.1140
- acc: 0.9622 - val_loss: 0.3792 - val_acc: 0.8820

Model Evaluation

```
In [32]: X_test = pad_sequences(sequences, maxlen=max_len)
Y_test = np.asarray(labels)

model.evaluate(X_test,Y_test)
```

25000/25000 [=====] - 139s 6ms/step

```
Out[32]: [2.317019919424057, 0.4965600073337555]
```

```
In [33]: history_dict = history.history


acc = history_dict['acc']
val_acc = history_dict['val_acc']
loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']
epochs = range(1, len(acc) + 1)

# Plotting metrics
plt.plot(epochs, acc, 'bo', label = 'Training accuracy')
plt.plot(epochs, val_acc, 'b', label = 'Validation accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()

plt.figure()

plt.plot(epochs, loss_values, 'bo', label = 'Training Loss')
plt.plot(epochs, val_loss_values, 'b', label = 'Validation loss')
plt.title('Training and Validation Loss')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()

plt.show()
```



Assignment 10.4

```
In [34]: from keras.models import Sequential
from keras import layers
from keras.optimizers import RMSprop
```

```
In [35]: model = Sequential()
model.add(Embedding(max_words,
                    embedding_dim,
                    input_length=max_len))
model.add(layers.Conv1D(32, 7, activation='relu'))
model.add(layers.MaxPooling1D(5))
model.add(layers.Conv1D(32, 7, activation='relu'))
model.add(layers.GlobalMaxPooling1D())
model.add(layers.Dense(1, activation='sigmoid'))
model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
embedding_3 (Embedding)	(None, 2234, 100)	1000000

conv1d_1 (Conv1D)	(None, 2228, 32)	22432

max_pooling1d_1 (MaxPooling1D)	(None, 445, 32)	0

conv1d_2 (Conv1D)	(None, 439, 32)	7200

global_max_pooling1d_1 (GlobalMaxPooling1D)	(None, 32)	0

dense_4 (Dense)	(None, 1)	33
=====		
Total params: 1,029,665		
Trainable params: 1,029,665		
Non-trainable params: 0		

```
In [36]: model.compile(optimizer=RMSprop(learning_rate=1e-4),
                    loss='binary_crossentropy',
                    metrics=['acc'])
history = model.fit(X_train, Y_train,
                    epochs=10,
                    batch_size=128,
                    validation_data=(X_val, Y_val))
```

C:\Users\bibek\anaconda3\envs\dsc650\lib\site-packages\tensorflow_core\python\framework\indexed_slices.py:433: UserWarning: Converting sparse IndexedSlices to a dense Tensor of unknown shape. This may consume a large amount of memory.

"Converting sparse IndexedSlices to a dense Tensor of unknown shape. "

Train on 24500 samples, validate on 500 samples

Epoch 1/10

24500/24500 [=====] - 92s 4ms/step - loss: 0.6922 - acc: 0.5345 - val_loss: 0.6905 - val_acc: 0.5860

Epoch 2/10

24500/24500 [=====] - 92s 4ms/step - loss: 0.6838 - acc: 0.6572 - val_loss: 0.6776 - val_acc: 0.7220

Epoch 3/10

24500/24500 [=====] - 96s 4ms/step - loss: 0.6540 - acc: 0.7644 - val_loss: 0.6225 - val_acc: 0.7580

Epoch 4/10

24500/24500 [=====] - 91s 4ms/step - loss: 0.5540 - acc: 0.8018 - val_loss: 0.5019 - val_acc: 0.7980

Epoch 5/10

24500/24500 [=====] - 95s 4ms/step - loss: 0.4177 - acc: 0.8419 - val_loss: 0.4100 - val_acc: 0.8180

Epoch 6/10

24500/24500 [=====] - 92s 4ms/step - loss: 0.3329 - acc: 0.8698 - val_loss: 0.3870 - val_acc: 0.8320

Epoch 7/10

24500/24500 [=====] - 95s 4ms/step - loss: 0.2885 - acc: 0.8853 - val_loss: 0.3501 - val_acc: 0.8620

Epoch 8/10

24500/24500 [=====] - 94s 4ms/step - loss: 0.2599 - acc: 0.8979 - val_loss: 0.3345 - val_acc: 0.8600

Epoch 9/10

24500/24500 [=====] - 96s 4ms/step - loss: 0.2382 - acc: 0.9073 - val_loss: 0.3457 - val_acc: 0.8620

Epoch 10/10

24500/24500 [=====] - 100s 4ms/step - loss: 0.2208 - acc: 0.9152 - val_loss: 0.3218 - val_acc: 0.8700

Model Evaluation


```
In [37]: X_test = pad_sequences(sequences, maxlen=max_len)
Y_test = np.asarray(labels)

model.evaluate(X_test,Y_test)
```

25000/25000 [=====] - 19s 746us/step

```
Out[37]: [1.7728079881000518, 0.4962399899959564]
```

```
In [38]: history_dict = history.history



acc = history_dict['acc']
val_acc = history_dict['val_acc']
loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']
epochs = range(1, len(acc) + 1)

# Plotting metrics
plt.plot(epochs, acc, 'bo', label = 'Training accuracy')
plt.plot(epochs, val_acc, 'b', label = 'Validation accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()

plt.figure()

plt.plot(epochs, loss_values, 'bo', label = 'Training Loss')
plt.plot(epochs, val_loss_values, 'b', label = 'Validation loss')
plt.title('Training and Validation Loss')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()

plt.show()
```



In []: